
Streams

Hungry? Why wait? (for the rest of your data)

Streams: Intro

"We should have some ways of connecting programs like garden hose--screw in another segment when it becomes necessary to massage data in another way. This is the way of IO also."

-- Doug McIlroy. October 11, 1964
Original developer of Unix pipes

Streams: Intro

Streams are like Arrays, but laid out in time,
rather than in memory.

Streams: Intro

Streams are useful for:

- Start processing data before the last bit is received
 - Start sending a response before the request is completely received
- Process data in pieces without buffering the whole
 - Especially useful for processing large files too big to fit in memory
- Normalized API for handling disparate data types

Streams: Types

There are 5 kinds of streams:

- readable
 - writable
 - transform
 - duplex
 - “classic”
-

Streams: Commonalities (.pipe())

All streams use *.pipe()* to pair inputs with outputs.

.pipe() connects a readable stream's output to a writable stream's input

```
1 // .pipe() chaining
2 a.pipe(b).pipe(c).pipe(d)
3
4 // Equivalent to
5
6 a.pipe(b)
7 b.pipe(c)
8 c.pipe(d)
```

Streams: Readable

A readable stream is a stream that outputs data.

Streams: Readable

A readable stream is a stream that **outputs** data.

Examples of readable streams:

- http responses, on the client
 - http requests, on the server
 - fs read streams
 - zlib streams
 - crypto streams
 - tcp sockets
 - child process stdout and stderr
 - process.stdin
-

Streams: Readable

Reading a file **without** streams:

```
1 let http = require('http')
2 let fs = require('fs')
3
4 let server = http.createServer(function (req, res) {
5   fs.readFile(__filename, function (err, data) {
6     res.end(data)
7   })
8 })
9 server.listen(8000)
```

Streams: Readable

Reading a file **without** streams:

Pros:

- Simple

Cons:

- Verbose
 - Must store a copy of *__filename* in memory for **every request**
 - Must wait for *__filename* to finish loading before sending response
-

Streams: Readable

Reading a file **with** streams:

```
1 let http = require('http')
2 let fs = require('fs')
3
4 let server = http.createServer(function (req, res) {
5   fs.createReadStream(__filename).pipe(res)
6 })
7 server.listen(8000)
```

Streams: Readable

Reading a file **with** streams:

Pros:

- Simpler code
- Concise
- Low latency
- Low memory pressure

Cons:

- More conceptually complex
-

Streams: Readable

Add compression. Just 1 line!

```
1 let http = require('http')
2 let fs = require('fs')
3 let oppressor = require('oppressor')
4
5 let server = http.createServer(function (req, res) {
6   fs.createReadStream(__filename)
7     .pipe(oppressor(req))
8     .pipe(res)
9 })
```

Streams: Readable

Readable streams output data that can be inputted into other stream types:

- writable
- transform
- duplex

```
1 readableStream.pipe(writableOrTransformOrDuplexStream)
```

Streams: Readable

Create an arbitrary readable stream:

```
1 let Readable = require('stream').Readable
2
3 let readableStream = new Readable
4 readableStream.push('hello ')
5 readableStream.push('world\n')
6 // Send null to end stream
7 readableStream.push(null)
8
9 readableStream.pipe(process.stdout)
```

Streams: Readable

Implement `._read()` to build a custom readable stream:

```
1 let Readable = require('stream').Readable
2 let readableStream = Readable()
3
4 let c = 97
5 readableStream._read = () => {
6   readableStream.push(String.fromCharCode(c++))
7   if (c > 'z'.charCodeAt(0)) readableStream.push(null)
8 }
9
10 readableStream.pipe(process.stdout)
```


Streams: Readable

Notes:

- Data can be pushed before or after being piped
 - Similar to promises holding their value
 - Different from promises because they can only be read once
- Readable streams start in **non-flowing** or paused mode
- *.pipe()* unpauses or “resumes” a stream
- Adding a *'data'* event listener will also resume a readable stream

Streams: Readable

Manually consume a readable stream by calling *.read()*:

```
1 process.stdin.on('readable', () => {  
2   let buffer = process.stdin.read()  
3   console.log(buffer)  
4 })
```

Streams: Writable

A writable stream is a stream that **inputs** data.

Examples of writable streams:

- http requests, on the client
 - http responses, on the server
 - fs write streams
 - zlib streams
 - crypto streams
 - tcp sockets
 - child process stdin
 - process.stdout
 - process.stderr
-

Streams: Writable

Writing a file **without** streams:

```
1 let http = require('http')
2 let fs = require('fs')
3
4 let server = http.createServer(function (req, res) {
5   let buffer
6   req.on('data', data => {
7     buffer += data
8   })
9   req.on('end', () => {
10     fs.writeFile(__dirname + '/data.txt', buffer, function (err, data) {
11       res.end()
12     })
13   })
14 })
15 server.listen(8000)
```

Streams: Writable

Writing a file **without** streams:

Pros:

- ??? (No longer simple)

Cons:

- Verbose
 - Must store a copy of *req.body* in memory for **every request**
 - Must wait for *request* to complete before writing file
-

Streams: Writable

Writing a file **with** streams:

```
1 let http = require('http')
2 let fs = require('fs')
3
4 let server = http.createServer(function (req, res) {
5   req.pipe(fs.createWriteStream(__dirname + '/data.txt'))
6 })
7 server.listen(8000)
```

Streams: Writable

Writing a file **with** streams:

Pros:

- Simpler code
- Concise
- Low latency
- Low memory pressure

Cons:

- ~~More conceptually complex~~
-

Streams: Writable

Writable streams accept input from other stream's output:

- readable
- transform
- duplex

```
1 readableOrTransformOrDuplexStream.pipe(writableStream)
```

Streams: Writable

Create a writable stream:

```
1 let fs = require('fs')
2 let writableStream = fs.createWriteStream(__filename)
3
4 writableStream.write('hello ')
5
6 setImmediate(() => writableStream.end('world\n'))
```

Streams: Writable

Implement `._write()` to build a custom writable stream:

```
1 let Writable = require('stream').Writable
2 let writableStream = new Writable
3
4 // Implement the _write function to consume data
5 writableStream._write = (chunk, enc, next) => {
6   console.log(chunk)
7   next()
8 }
9
10 process.stdin.pipe(writableStream)
```

Streams: Writable

Notes:

- Data can be pushed before or after being piped
 - Similar to promises holding their value
 - Different from promises because they can only be read once
- Readable streams start in **non-flowing** or paused mode
- *.pipe()* unpauses or “resumes” a stream
- Adding a *'data'* event listener will also resume a readable stream

Streams: Writable

Manually consume a readable stream by calling *.read()*:

```
1 process.stdin.on('readable', () => {  
2   let buffer = process.stdin.read()  
3   console.log(buffer)  
4 })
```

Streams: Transform

A transform (aka “through”) stream both **inputs** and **outputs** data.

A transform stream typically transforms data in some way.

```
1 let through = require('through')
2 process.stdin.pipe(through(console.log, ()=>console.log('\n')))
```

Transform stream examples: zlib streams, crypto streams

Streams: Duplex

A duplex stream both inputs and outputs data.

A duplex stream is a combination readable and writable stream.

(Implements both `._read()` and `._write()`)

```
1 let fs = require('fs')
2 let request = require('request')
3 let duplexStream = request('http://google.com')
4 fs.createReadStream(__filename).pipe(duplexStream).pipe(process.stdout)
```

Duplex stream examples: tcp sockets, zlib streams, crypto streams

Streams: Classic or Legacy

- Classic streams are *EventEmitters* with no buffering
 - aka “push streams”
 - Similar API
 - Always in “flowing” mode
 - Start out unpaused, which means you can drop ‘data’ events
 - Adding a ‘data’ event listener to a stream2 will cause it to fall back into legacy mode
 - No special `._read()` or `._write()` support
 - Must handle buffering and back-pressure manually
-